Biniyam Berga

Data Scientist

Machine Learning, Analytics & LLM Evaluation

biniyamfikiru4@gmail.com
+251-925-920458
Addis Ababa, Ethiopia

linkedin.com/in/biniyam-berga github.com/biniyam44 biniyam44.github.io

Data Scientist with 4+ years of experience in machine learning, analytics, and LLM evaluation. Builds predictive models (XGBoost, Random Forest, regression, clustering), SQL-based forecasting models, and Power BI/Tableau dashboards for business stakeholders, alongside Python evaluation pipelines and FastAPI/Docker tooling. Currently a Subject Matter Expert evaluating frontier AI models at Pareto AI; previously promoted to Lead Python Developer at Turing, leading an 8-person team producing LLM training data and evaluation suites.

CORE SKILLS

AI/LLM Evaluation: RLHF, prompt engineering, rubric design, red-teaming & adversarial testing, jailbreak/prompt-injection testing, multimodal (image/text) annotation & QA

Data Science & ML: forecasting, regression & clustering, XGBoost, Random Forest, PyTorch, TensorFlow, evaluation pipelines

Programming & Infrastructure: Python, SQL, PostgreSQL, FastAPI, Docker, Git, NumPy, pandas, OpenCV

Analytics & BI: Power BI, Tableau, SSRS, Excel/Google Sheets, statistical analysis

PROFESSIONAL EXPERIENCE

Pareto AI (Contract) — Oct 2025 – Present

Subject Matter Expert, Data Scientist

Skills: Python, SQL, Excel, spreadsheets, machine learning, DevOps, GitHub, Docker.

- Evaluate frontier AI model outputs across text, code, and multimodal tasks for reasoning quality, factual accuracy, and instruction alignment, feeding structured scores into RLHF training pipelines.
- Design and stress-test prompts to surface model failure modes, informing rubric revisions used across the evaluation team.
- Build and validate forecasting and structured evaluation models in Python and Excel, keeping results reproducible across reviewers.
- Hired directly at the Subject Matter Expert tier for data science evaluation work, one level above standard annotator roles.

Turing (Contract) — Sep 2023 – Mar 2026

Lead Python Developer (promoted from Python Developer) — **Aug 2024 – Mar 2026**

Skills: Python, SQL, Excel, spreadsheets, machine learning, DevOps, GitHub, Docker.

- Lead an 8-person team of trainers and reviewers producing human-in-the-loop training data and code-evaluation pipelines for LLM fine-tuning.
- Designed QA rubrics and checks to catch hallucinations, hidden reasoning steps, and data leakage; calibrated reviewer scoring across the 8-person team and ran onboarding on annotation standards and anti-cheating detection.
- Built evaluation suites benchmarking model reasoning, safety, and robustness; shipped model-serving and tooling APIs with FastAPI, Docker, and Postgres to make workflows reproducible.
- Worked on video labeling and annotation and image labeling and annotation alongside coding and multimodal evaluation.

Python Developer **Sep 2023 – Jul 2024**

- Trained LLMs to complete incomplete Python repositories by parsing READMEs, test suites, and dependent libraries, enforcing original-implementation rules that blocked copied solutions.
- Designed and enforced rubrics and QA checks to detect hallucinations, hidden reasoning steps, and data leakage; built feedback loops that raised dataset quality.
- Performed prompt engineering and developed evaluation suites to benchmark model reasoning, safety, and robustness.

IoToad (Contract) **Dec 2022 – Dec 2023**

Data Analyst

- Built dashboard standards with SSRS and custom Tableau/Power BI dashboards, cutting data retrieval time and giving stakeholders real-time visibility across departments.
- Directed SQL-based forecasting and budgeting models, tightening budget allocation accuracy and reducing financial discrepancies.
- Integrated data sources via SQL Server and cleaned large datasets with SQL/Power BI, improving data quality and processing speed for the analytics team.

Camara Education Ethiopia **Jun 2022 – Oct 2022**

Data Analyst

- Built predictive models with XGBoost and Random Forest and applied regression/clustering techniques to operational data, improving forecasting accuracy and customer targeting.
- Established data pipelines with Apache NiFi and designed Tableau/Power BI dashboards, speeding up ingestion and giving stakeholders faster access to insights.
- Used SQL for exploratory analysis and feature engineering to support model development.

PROJECTS

AfterQuery Project collaborations

AI engineering and evaluation

Tools: Python, SQL, Ruby, C++, Docker, GitHub.

- Contribute to Terminal Bench, Long Horizon Terminal Bench, and MLEBench projects.
- Develop repository-based coding tasks, behavioral and regression tests, and reference solutions.
- Investigate failures and validate reproducible Docker environments and test harnesses.

AI Red-Teaming & Adversarial Testing

- Designed multi-step prompt attack sequences to probe LLM guardrails, reporting reproducible exploits with proposed mitigations.
- Simulated scenarios where model outputs could leak sensitive schema or training data, building detection heuristics to catch these failures.
- Systematically induced reasoning failures using contradictory constraints and chained dependencies to stress-test model robustness.

EDUCATION

Adama Science and Technology University | BSc, Computer Science & Engineering and Electronics & Communication Engineering (Dual Major) | **July 2023**

St. Mary's University | BA, Management | **February 2023**

ALX | Certificates: Software Engineering (Back-End Specialty) · Business Intelligence & Data Analyst · Financial Modelling & Valuation Analyst